

Use Machine Learning to detect Age-Related conditions

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.express as px
%matplotlib inline
import seaborn as sns
import gc
import re as re
from collections import Counter

from tqdm.auto import tqdm
import math
from sklearn.model_selection import KFold, StratifiedKFold, train_test_split, GridSearchCV
from sklearn.metrics import roc_auc_score, accuracy_score, confusion_matrix, ConfusionMatrixDisplay, RocCurveDisplay
from sklearn.preprocessing import LabelEncoder, StandardScaler

import warnings
warnings.filterwarnings('ignore')

from glob import glob
from pathlib import Path
import joblib
import pickle
import os
import random

import time
from xgboost import XGBClassifier
from sklearn.metrics import log_loss
%matplotlib inline

/opt/conda/lib/python3.10/site-packages/scipy/_init__.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}"
```

Load Dataset

```
In [2]: train      = pd.read_csv('/input/icr-identify-age-related-conditions/train.csv')
test       = pd.read_csv('/input/icr-identify-age-related-conditions/test.csv')

In [3]: train.head()

Out[3]:    Id   AB    AF    AH    AM    AR    AX    AY    AZ    BC ...    FL    FR    FS    GB
0  000ff2bfdf9  0.209377  3109.03329  85.200147  22.394407  8.138688  0.699861  0.025578  9.812214  5.555634 ...  7.298162  1.73855  0.094822  11.339138  7e
1  007255e47698  0.145282  978.76416  85.200147  36.968889  8.138688  3.632190  0.025578  13.517790  1.229900 ...  0.173229  0.49706  0.568932  9.292698  7e
2  013f2bd269f5  0.470030  2635.10654  85.200147  32.360553  8.138688  6.732840  0.025578  12.824570  1.229900 ...  7.709560  0.97556  1.198821  37.077772  8e
3  043ac50845d5  0.252107  3819.65177  120.201618  77.112203  8.138688  3.685344  0.025578  11.053708  1.229900 ...  6.122162  0.49706  0.284466  18.529584  8e
4  044fb8a146ec  0.380297  3733.04844  85.200147  14.103738  8.138688  3.942255  0.054810  3.396778  102.151980 ...  8.153058  48.50134  0.121914  16.408728  14e
```

5 rows × 58 columns

EDA

```
In [4]: # summary table function
def summary(df):
    print(f'data shape: {df.shape}')
    summ = pd.DataFrame(df.dtypes, columns=['data type'])
    summ['#missing'] = df.isnull().sum().values
    summ['%missing'] = df.isnull().sum().values / len(df)* 100
    summ['#unique'] = df.nunique().values
    desc = pd.DataFrame(df.describe(include='all').transpose())
    summ['min'] = desc['min'].values
    summ['max'] = desc['max'].values
    summ['first quartile'] = desc.loc[:, '25%'].values
    summ['second quartile'] = desc.loc[:, '50%'].values
    summ['third quartile'] = desc.loc[:, '75%'].values

    return summ
```

```
In [5]: summary(train)
data shape: (617, 58)
```

Out[5]:

	data type	#missing	%missing	#unique	min	max	first quartile	second quartile	third quartile
Id	object	0	0.000000	617	NaN	NaN	NaN	NaN	NaN
AB	float64	0	0.000000	217	0.081187	6.161666	0.252107	0.354659	0.559763
AF	float64	0	0.000000	599	192.59328	28688.18766	2197.34548	3120.31896	4361.63739
AH	float64	0	0.000000	227	85.200147	1910.123198	85.200147	85.200147	113.73954
AM	float64	0	0.000000	605	3.177522	630.51823	12.270314	20.53311	39.139886
AR	float64	0	0.000000	130	8.138688	178.943634	8.138688	8.138688	8.138688
AX	float64	0	0.000000	427	0.699861	38.27088	4.128294	5.031912	6.431634
AY	float64	0	0.000000	148	0.025578	10.315851	0.025578	0.025578	0.036845
AZ	float64	0	0.000000	484	3.396778	38.971568	8.12958	10.46132	12.969516
BC	float64	0	0.000000	259	1.2299	1463.693448	1.2299	1.2299	5.081244
BD	float64	0	0.000000	617	1693.62432	53060.59924	4155.70287	4997.96073	6035.8857
BN	float64	0	0.000000	53	9.8868	29.3073	19.4205	21.186	23.6577
BP	float64	0	0.000000	612	72.948951	2447.81055	156.847239	193.908816	247.803462
BQ	float64	60	9.724473	515	1.331155	344.644105	27.834425	61.642115	134.009015
BR	float64	0	0.000000	566	51.216883	179250.2529	424.990642	627.417402	975.649259
BZ	float64	0	0.000000	115	257.432377	50092.4593	257.432377	257.432377	257.432377
CB	float64	2	0.324149	553	12.49976	2271.436167	23.317567	42.55433	77.310097
CC	float64	3	0.486224	602	0.176874	4.103032	0.563688	0.658715	0.772206
CD	float64	0	0.000000	584	23.3876	633.534408	64.724192	79.819104	99.81352
CF	float64	0	0.000000	586	0.510888	200.967526	5.066306	9.123	13.565901
CH	float64	0	0.000000	135	0.003184	0.224074	0.023482	0.02786	0.034427
CL	float64	0	0.000000	123	1.050225	31.688153	1.050225	1.050225	1.228445
CR	float64	0	0.000000	595	0.069225	3.039675	0.589575	0.7308	0.85935
CS	float64	0	0.000000	576	13.784111	267.942823	29.782467	34.83513	40.529401
CU	float64	0	0.000000	307	0.137925	4.951507	1.070298	1.351665	1.660617
CW	float64	0	0.000000	426	7.03064	64.521624	7.03064	36.019104	37.935832
DA	float64	0	0.000000	611	6.9064	210.33092	37.94252	49.18094	61.40876
DE	float64	0	0.000000	616	35.998895	2103.40519	188.81569	307.509595	507.8962
DF	float64	0	0.000000	137	0.23868	37.895013	0.23868	0.23868	0.23868
DH	float64	0	0.000000	191	0.040995	1.060404	0.295164	0.358023	0.426348
DI	float64	0	0.000000	571	60.23247	1049.168078	102.703553	130.05063	165.836955
DL	float64	0	0.000000	604	10.3456	326.2362	78.23224	96.26496	110.64068
DN	float64	0	0.000000	576	6.339496	62.808096	20.888264	25.2488	30.544224
DU	float64	1	0.162075	253	0.005518	161.355315	0.005518	0.251741	1.05869
DV	float64	0	0.000000	39	1.74307	25.19293	1.74307	1.74307	1.74307
DY	float64	0	0.000000	590	0.804068	152.355164	14.715792	21.642456	34.058344
EB	float64	0	0.000000	439	4.926396	94.95858	5.965392	8.149404	10.503048
EE	float64	0	0.000000	513	0.286201	18.324926	1.648679	2.616119	3.91007
EG	float64	0	0.000000	610	185.5941	30243.75878	1111.160625	1493.817413	1905.701475
EH	float64	0	0.000000	127	0.003042	42.569748	0.003042	0.085176	0.237276
EJ	object	0	0.000000	2	NaN	NaN	NaN	NaN	NaN
EL	float64	60	9.724473	311	5.394675	109.125159	30.927468	71.949306	109.125159
EP	float64	0	0.000000	275	78.526968	1063.594578	78.526968	78.526968	112.766654
EU	float64	0	0.000000	455	3.828384	6501.264448	4.324656	22.641144	49.085352
FC	float64	1	0.162075	600	7.534128	3030.655824	25.815384	36.394008	56.714448
FD	float64	0	0.000000	337	0.29685	1578.654237	0.29685	1.870155	4.880214
FE	float64	0	0.000000	615	1563.136688	143224.6823	5164.66626	7345.143424	10647.95165
FI	float64	0	0.000000	498	3.58345	35.851039	8.523098	9.945452	11.516657
FL	float64	1	0.162075	388	0.173229	137.932739	0.173229	3.028141	6.238814
FR	float64	0	0.000000	435	0.49706	1244.22702	0.49706	1.131	1.51206
FS	float64	2	0.324149	161	0.06773	31.365763	0.06773	0.250601	0.535067

	data type	#missing	%missing	#unique	min	max	first quartile	second quartile	third quartile
GB	float64	0	0.000000	560	4.102182	135.781294	14.036718	18.771436	25.608406
GE	float64	0	0.000000	264	72.611063	1497.351958	72.611063	72.611063	127.591671
GF	float64	0	0.000000	611	13.038894	143790.0712	2798.992584	7838.27361	19035.70924
GH	float64	0	0.000000	596	9.432735	81.210825	25.034888	30.608946	36.863947
GI	float64	0	0.000000	615	0.897628	191.194764	23.011684	41.007968	67.931664
GL	float64	1	0.162075	355	0.001129	21.978	0.124392	0.337827	21.978
Class	int64	0	0.000000	2	0.0	1.0	0.0	0.0	0.0

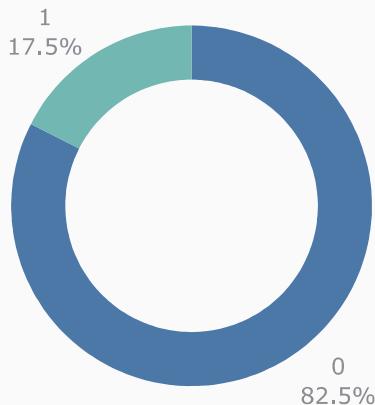
- There are 56 X variables and 1 target(y) variable, while 1 variable(id) is not relevant for modelling
- All variable types are float64, except for the EJ variable which is of type object.
- 9 variables have missing value. The missing value ratio is not very high.

```
In [6]: # select numerical and categorical variables respectively.
num_cols = test.select_dtypes(include=['float64']).columns.tolist()
cat_cols = test.select_dtypes(include=['object']).columns.tolist()
cat_cols.remove('Id')
```

Inspecting the class imbalance

```
In [7]: fig2 = px.pie(train, names='Class',
                  height=400, width=600,
                  hole=0.7,
                  title='target class Overview',
                  color_discrete_sequence=['#4c78a8', '#72b7b2'])
fig2.update_traces(hovertemplate=None, textposition='outside', textinfo='percent+label', rotation=0)
fig2.update_layout(margin=dict(t=100, b=30, l=0, r=0), showlegend=False,
                   plot_bgcolor="#fafafa", paper_bgcolor="#fafafa",
                   title_font=dict(size=20, color="#555", family="Lato, sans-serif"),
                   font=dict(size=17, color="#8a8d93"),
                   hoverlabel=dict(bgcolor="#444", font_size=13, font_family="Lato, sans-serif"))
fig2.show()
```

target class Overview



The dataset is quite imbalanced. Maybe we need to try a few over-sampling methods like SMOTE in future versions.

Distribution of numerical variables

```
In [8]: figsize = (4*4, 20)
fig = plt.figure(figsize=figsize)
for idx, col in enumerate(num_cols):
    ax = plt.subplot(11,5, idx + 1)
    sns.kdeplot(
        data=train, hue='Class', fill=True,
        x=col, palette=[ '#9E3F00', 'red'], legend=False
    )

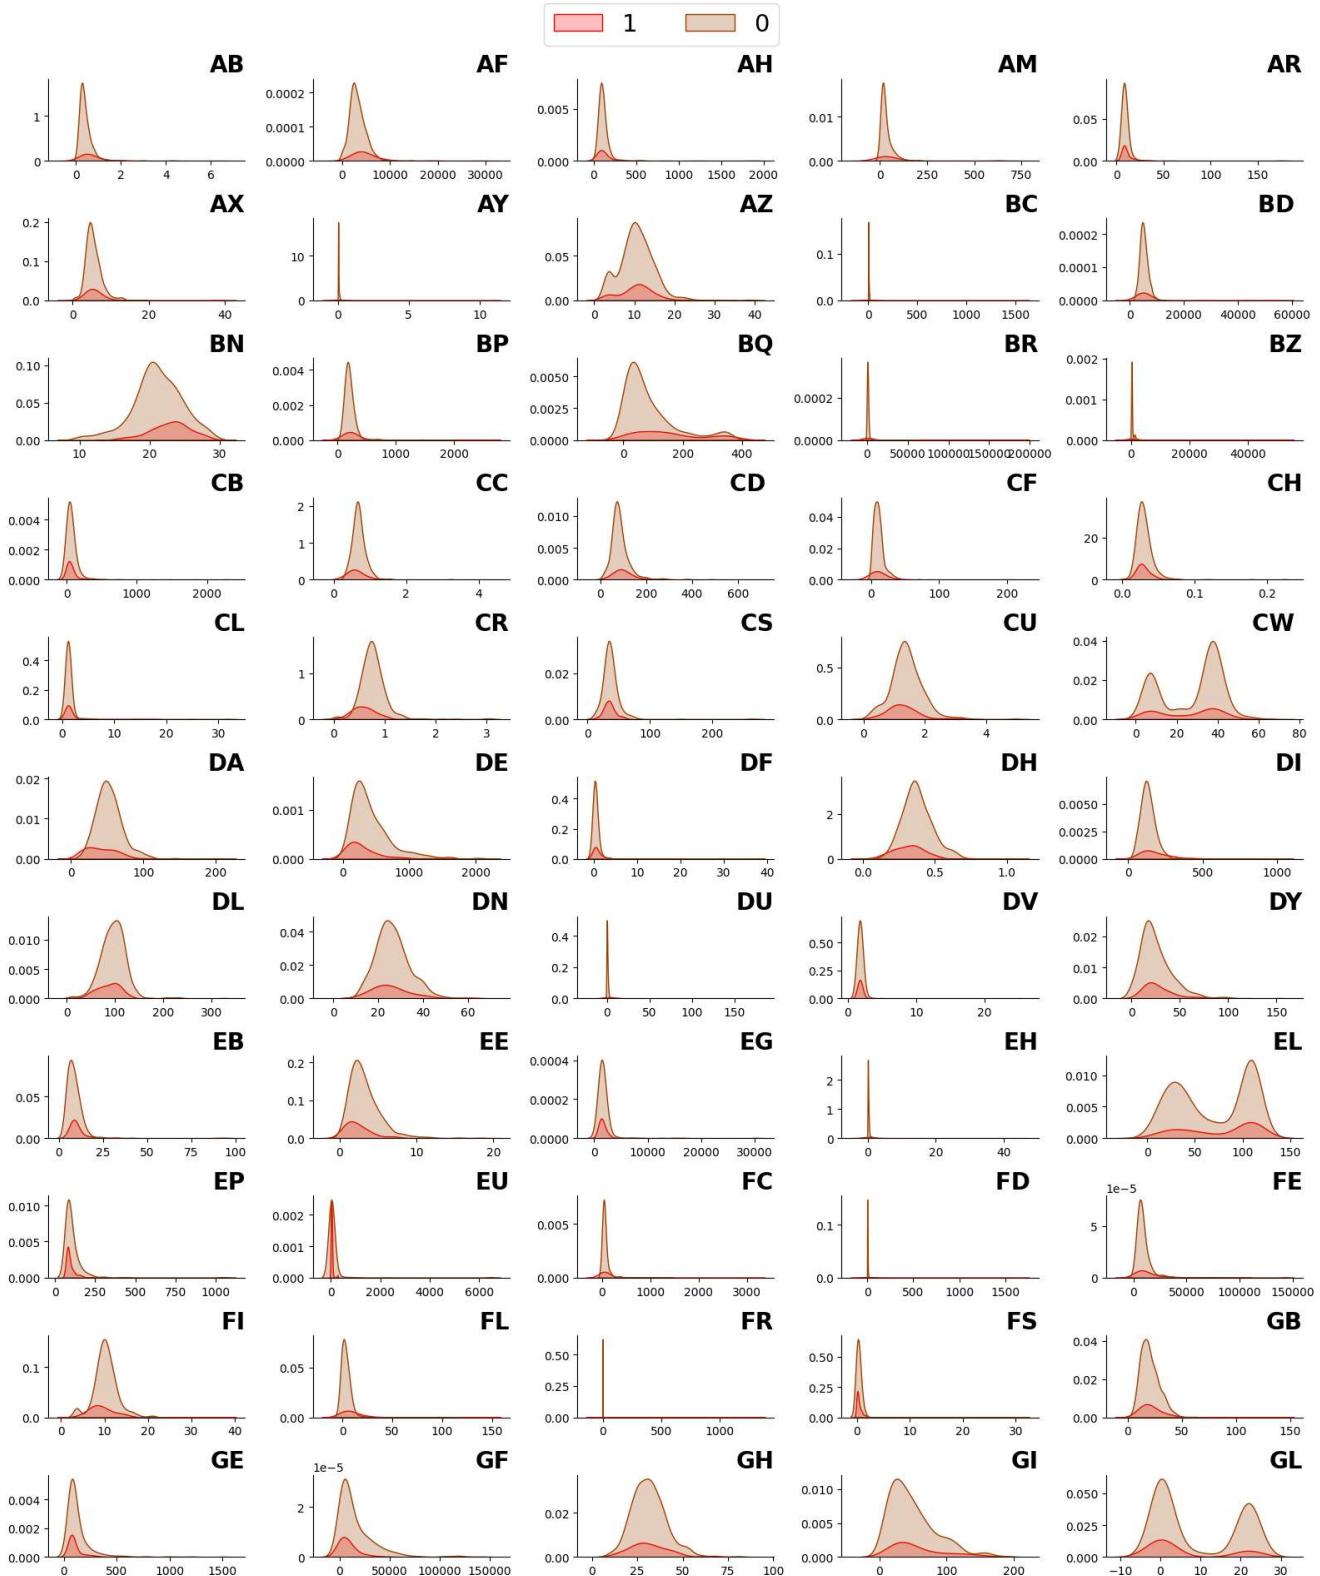
    ax.set_ylabel(''); ax.spines['top'].set_visible(False),
    ax.set_xlabel(''); ax.spines['right'].set_visible(False)
    ax.set_title(f'{col}', loc='right',
                weight='bold', fontsize=20)
```

```

fig.suptitle('Features vs Target\n\n\n', ha='center', fontweight='bold', fontsize=21)
fig.legend([1, 0], loc='upper center', bbox_to_anchor=(0.5, 0.96), fontsize=21, ncol=3)
plt.tight_layout()
plt.show()

```

Features vs Target



"The distribution of most numerical variables differs, indicating that these variables possess predictive power to some extent."

"However, scaling might be necessary as they correspond to different scales."

Correlation heatmap for numerical variables

```

In [9]: def plot_top_correlations(df: pd.core.frame.DataFrame, n: int, title_name: str='Top Correlations') -> None:
    # Calculate correlation between all variables

```

```

corr = df.corr()

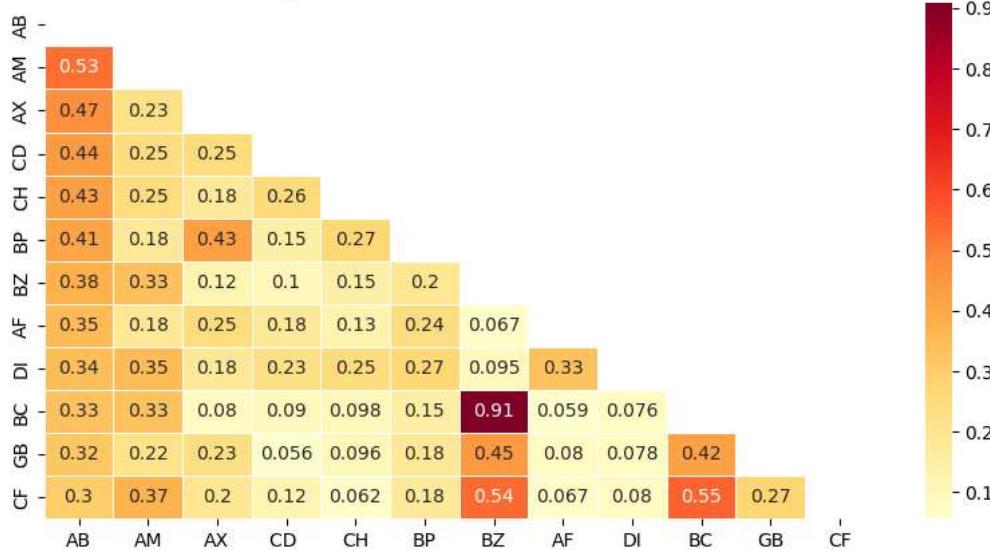
# Select variables having highest absolute correlation
top_corr_cols = corr.abs().nlargest(n, columns=corr.columns).index
top_corr = corr.loc[top_corr_cols, top_corr_cols]

fig, axes = plt.subplots(figsize=(10, 5))
mask = np.zeros_like(top_corr)
mask[np.triu_indices_from(mask)] = True
sns.heatmap(top_corr, mask=mask, linewidths=.5, cmap='YlOrRd', annot=True)
plt.title('title_name')
plt.show()

# Plot heatmap of top 12 correlations in training data
plot_top_correlations(train[num_cols], 12, 'Top 12 Correlations in Train Dataset')

```

Top 12 Correlations in Train Dataset



"Among the top 12 highly correlated numerical variables, while some exhibit strong correlations, the overall correlation among these health-related data variables is not very high."

Model training using XGB, LightGBM, CatBoost

```
In [10]: import lightgbm as lgb
import xgboost as xgb
from catboost import Pool, CatBoostRegressor, CatBoostClassifier
```

```
In [11]: class CFG:
    VER = 1
    AUTHOR = 'maverick'
    COMPETITION = 'icr-identify-age-related-conditions'
    DATA_PATH = Path('/input/icr-identify-age-related-conditions')
    OOF_DATA_PATH = Path('./oof')
    MODEL_DATA_PATH = Path('./models')
    METHOD_LIST = ['lightgbm', 'xgboost', 'catboost']
    seed = 3407 #52
    n_folds = 10 #replaced 20
    target_col = 'Class'
    metric = 'balanced_log_loss'
    metric_maximize_flag = False
    num_boost_round = 50500
    early_stopping_round = 500
    verbose = 2000
    boosting_type = 'dart'
    lgb_params = {
        'objective': 'binary', # 'binary', 'multiclass'
        'metric': None, # 'auc', 'multi_logloss'
        # 'num_class': None,
        'boosting': boosting_type,
        'device_type':'gpu',
        'learning_rate': 0.005,
        'num_leaves': 5,
        'feature_fraction': 0.50,
        'bagging_fraction': 0.80,
        'lambda_l1': 2,
        'lambda_l2': 4,
        'n_jobs': -1,
        'is_unbalance':True, #added balancing
        'verbose': -1, #added silence
        # 'min_data_in_leaf': 40,
        # 'bagging_freq': 10,
        'seed': seed,
```

```

}
xgb_params = {
    'objective': 'binary:logistic',
    'eval_metric': 'logloss',
    'learning_rate': 0.005,
    'max_depth': 4,
    'colsample_bytree': 0.50,
    'subsample': 0.80,
    'eta': 0.03,
    'gamma': 1.5,
    # 'Lambda': 70,
    # 'min_child_weight': 8,
    # 'eval_metric':'logloss',
    # 'tree_method': 'gpu_hist',
    # 'predictor': 'gpu_predictor',
    'random_state': seed,
}

cat_params = {
    'learning_rate': 0.005,
    'iterations': num_boost_round,
    'depth': 4, #
    'colsample_bylevel': 0.50,
    'subsample': 0.80,
    'l2_leaf_reg': 3, # 3, 30
    'random_seed': seed,
    'auto_class_weights': 'Balanced'
    # 'task_type': 'GPU',
}

```

In [12]: !mkdir oof
!mkdir models

In [13]: test[CFG.target_col] = -1
all_df = pd.concat([train, test])

Setting Seed

In [14]: def seed_everything(seed):
 random.seed(seed)
 np.random.seed(seed)
 os.environ['PYTHONHASHSEED'] = str(seed)

Metric Calculation

In [15]: def balanced_log_loss(y_true, y_pred):
 y_pred = np.clip(y_pred, 1e-15, 1-1e-15)
 nc = np.bincount(y_true)
 w0, w1 = 1/(nc[0]/y_true.shape[0]), 1/(nc[1]/y_true.shape[0])
 balanced_log_loss_score = (-w0/nc[0]*(np.sum(np.where(y_true==0, 1, 0) * np.log(1-y_pred))) - w1/nc[1]*(np.sum(np.where(y_true!=0, 1, 0) * np.log(1-y_pred))))
 return balanced_log_loss_score

=====
LightGBM Metric
=====
def lgb_metric(y_pred, y_true):
 y_true = y_true.get_label()
 return 'balanced_log_loss', balanced_log_loss(y_true, y_pred), CFG.metric_maximize_flag

def xgb_metric(y_pred, y_true):
 y_true = y_true.get_label()
 return 'balanced_log_loss', balanced_log_loss(y_true, y_pred)

=====
Catboost Metric
=====
class CatboostMetric(object):
 def get_final_error(self, error, weight): return error
 def is_max_optimal(self): return CFG.metric_maximize_flag
 def evaluate(self, approxes, target, weight):
 error = balanced_log_loss(np.array(target), approxes)
 return error, 0

In [16]: def calc_log_loss_weight(y_true):
 nc = np.bincount(y_true)
 w0, w1 = 1/(nc[0]/y_true.shape[0]), 1/(nc[1]/y_true.shape[0])
 return w0, w1

def lightgbm_training(x_train: pd.DataFrame, y_train: pd.DataFrame, x_valid: pd.DataFrame, y_valid: pd.DataFrame, features: list, categorical_features: list):
 train_w0, train_w1 = calc_log_loss_weight(y_train)
 valid_w0, valid_w1 = calc_log_loss_weight(y_valid)
 lgb_train = lgb.Dataset(x_train, y_train, weight=y_train.map({0: train_w0, 1: train_w1}), categorical_feature=categorical_features)
 lgb_valid = lgb.Dataset(x_valid, y_valid, weight=y_valid.map({0: valid_w0, 1: valid_w1}), categorical_feature=categorical_features)
 model = lgb.train(
 params = CFG.lgb_params,
 train_set = lgb_train,
 num_boost_round = CFG.num_boost_round,

```

        valid_sets = [lgb_train, lgb_valid],
        early_stopping_rounds = CFG.early_stopping_round,
        verbose_eval = CFG.verbose,
        # feval = lgb_metric,
    )
# Predict validation
valid_pred = model.predict(x_valid)
return model, valid_pred
def xgboost_training(x_train: pd.DataFrame, y_train: pd.DataFrame, x_valid: pd.DataFrame, y_valid: pd.DataFrame, features: list, categorical_features: list):
    train_w0, train_w1 = calc_log_loss_weight(y_train)
    valid_w0, valid_w1 = calc_log_loss_weight(y_valid)
    xgb_train = xgb.DMatrix(data=x_train, label=y_train, weight=y_train.map({0: train_w0, 1: train_w1}))
    xgb_valid = xgb.DMatrix(data=x_valid, label=y_valid, weight=y_valid.map({0: valid_w0, 1: valid_w1}))
    model = xgb.train(
        CFG.xgb_params,
        dtrain = xgb_train,
        num_boost_round = CFG.num_boost_round,
        evals = [(xgb_train, 'train'), (xgb_valid, 'eval')],
        early_stopping_rounds = CFG.early_stopping_round,
        verbose_eval = CFG.verbose,
        # feval = xgb_metric,
        # maximize = CFG.metric_maximize_flag,
    )
# Predict validation
valid_pred = model.predict(xgb.DMatrix(x_valid), iteration_range=(0, model.best_ntree_limit))
return model, valid_pred

def catboost_training(x_train: pd.DataFrame, y_train: pd.DataFrame, x_valid: pd.DataFrame, y_valid: pd.DataFrame, features: list, categorical_features: list):
    train_w0, train_w1 = calc_log_loss_weight(y_train)
    valid_w0, valid_w1 = calc_log_loss_weight(y_valid)
    cat_train = Pool(data=x_train, label=y_train, weight=y_train.map({0: train_w0, 1: train_w1}), cat_features=categorical_features)
    cat_valid = Pool(data=x_valid, label=y_valid, weight=y_valid.map({0: valid_w0, 1: valid_w1}), cat_features=categorical_features)
    model = CatBoostClassifier(**CFG.cat_params) # , eval_metric = CatboostMetric
    model.fit(cat_train,
              eval_set=[cat_valid],
              early_stopping_rounds=CFG.early_stopping_round,
              verbose=CFG.verbose,
              use_best_model=True)
# Predict validation
valid_pred = model.predict_proba(x_valid)[:, 1]
return model, valid_pred

def gradient_boosting_model_cv_training(method: str, train_df: pd.DataFrame, features: list, categorical_features: list):
    # Create a numpy array to store out of folds predictions
    oof_predictions = np.zeros(len(train_df))
    oof_fold = np.zeros(len(train_df))
    kfold = StratifiedKFold(n_splits = CFG.n_folds, shuffle = True, random_state = CFG.seed)
    for fold, (train_index, valid_index) in enumerate(kfold.split(train_df, train_df[CFG.target_col])):
        print(f'-'*50)
        print(f'{method} training fold {fold + 1}')
        x_train = train_df[features].iloc[train_index]
        y_train = train_df[CFG.target_col].iloc[train_index]
        x_valid = train_df[features].iloc[valid_index]
        y_valid = train_df[CFG.target_col].iloc[valid_index]
        if method == 'lightgbm':
            model, valid_pred = lightgbm_training(x_train, y_train, x_valid, y_valid, features, categorical_features)
        if method == 'xgboost':
            model, valid_pred = xgboost_training(x_train, y_train, x_valid, y_valid, features, categorical_features)
        if method == 'catboost':
            model, valid_pred = catboost_training(x_train, y_train, x_valid, y_valid, features, categorical_features)

        # Save best model
        pickle.dump(model, open(CFG.MODEL_DATA_PATH / f'{method}_fold{fold + 1}_seed{CFG.seed}_ver{CFG.VER}.pkl', 'wb'))
        # Add to out of folds array
        oof_predictions[valid_index] = valid_pred
        oof_fold[valid_index] = fold + 1
    del x_train, x_valid, y_train, y_valid, model, valid_pred
    gc.collect()

    # Compute out of folds metric
    score = balanced_log_loss(train_df[CFG.target_col], oof_predictions)
    print(f'{method} our out of folds CV score is {score}')
    # Create a dataframe to store out of folds predictions
    oof_df = pd.DataFrame({'Id': train_df['Id'], CFG.target_col: train_df[CFG.target_col], f'{method}_prediction': oof_predictions, 'fold': oof_fold})
    oof_df.to_csv(CFG.MODEL_DATA_PATH / f'oof_{method}_seed{CFG.seed}_ver{CFG.VER}.csv', index = False)

```

Feature Selection

```
In [17]: numerical_features = ['AB', 'AF', 'AH', 'AM', 'AR', 'AX', 'AY', 'AZ', '#BC',
                                'BD', 'BN', 'BP', 'BQ', 'BR', 'BZ',
                                'CB', 'CC', 'CD', 'CF', 'CH', '#CL',
                                'CR', 'CS', 'CU', 'CW',
                                'DA', 'DE', 'DH', 'DI', 'DL', 'DN', 'DU', 'DV', 'DY',
                                'EB', 'EE', 'EG', 'EH', 'EL', 'EP', 'EU',
                                'FC', 'FD', 'FE', 'FI', 'FL', 'FR', 'FS',
                                'GB', 'GE', 'GF', 'GH', 'GI', 'GL']
categorical_features = ['EJ']
features = numerical_features + categorical_features
```

```
str2int_dict = {}
str2int_dict['EJ'] = {'A': 1, 'B': 0}
seed_everything(CFG.seed)
```

Preprocessing

```
In [18]: def Preprocessing(input_df: pd.DataFrame) -> pd.DataFrame:
    input_df = input_df.rename(columns={'BD': 'CD', 'CD': 'CW', 'CW': 'FD', 'FD': 'BD'}) #?
    output_df = input_df.copy()
    sc = StandardScaler()
    output_df[numerical_features] = sc.fit_transform(output_df[numerical_features]) #added scaling
    for col in categorical_features:
        output_df[col] = input_df[col].map(str2int_dict[col])
    return output_df
all_df = Preprocessing(all_df)

In [19]: train_df = all_df[all_df[CFG.target_col] != -1].copy()
test_df = all_df[all_df[CFG.target_col] == -1].copy()

In [20]: for method in CFG.METHOD_LIST:
    gradient_boosting_model_cv_training(method, train_df, features, categorical_features)
```

lightgbm training fold 1

[2000] training's binary_logloss: 0.240573 valid_1's binary_logloss: 0.380602
[4000] training's binary_logloss: 0.133936 valid_1's binary_logloss: 0.335777
[6000] training's binary_logloss: 0.083121 valid_1's binary_logloss: 0.313393
[8000] training's binary_logloss: 0.0549198 valid_1's binary_logloss: 0.311155
[10000] training's binary_logloss: 0.0411957 valid_1's binary_logloss: 0.323622
[12000] training's binary_logloss: 0.0319705 valid_1's binary_logloss: 0.333705
[14000] training's binary_logloss: 0.0269473 valid_1's binary_logloss: 0.343834
[16000] training's binary_logloss: 0.0230827 valid_1's binary_logloss: 0.352815
[18000] training's binary_logloss: 0.020925 valid_1's binary_logloss: 0.358794
[20000] training's binary_logloss: 0.0193932 valid_1's binary_logloss: 0.36387
[22000] training's binary_logloss: 0.0181028 valid_1's binary_logloss: 0.365714
[24000] training's binary_logloss: 0.0172531 valid_1's binary_logloss: 0.36771
[26000] training's binary_logloss: 0.0165144 valid_1's binary_logloss: 0.36882
[28000] training's binary_logloss: 0.016043 valid_1's binary_logloss: 0.370757
[30000] training's binary_logloss: 0.0158354 valid_1's binary_logloss: 0.372224
[32000] training's binary_logloss: 0.0156357 valid_1's binary_logloss: 0.375262
[34000] training's binary_logloss: 0.0154769 valid_1's binary_logloss: 0.377949
[36000] training's binary_logloss: 0.0153588 valid_1's binary_logloss: 0.379529
[38000] training's binary_logloss: 0.0152481 valid_1's binary_logloss: 0.381004
[40000] training's binary_logloss: 0.015148 valid_1's binary_logloss: 0.382658
[42000] training's binary_logloss: 0.015066 valid_1's binary_logloss: 0.383445
[44000] training's binary_logloss: 0.0149789 valid_1's binary_logloss: 0.384562
[46000] training's binary_logloss: 0.0149165 valid_1's binary_logloss: 0.384312
[48000] training's binary_logloss: 0.0148306 valid_1's binary_logloss: 0.384465
[50000] training's binary_logloss: 0.0147753 valid_1's binary_logloss: 0.384568

lightgbm training fold 2

[2000] training's binary_logloss: 0.255251 valid_1's binary_logloss: 0.305696
[4000] training's binary_logloss: 0.144125 valid_1's binary_logloss: 0.203399
[6000] training's binary_logloss: 0.091037 valid_1's binary_logloss: 0.159946
[8000] training's binary_logloss: 0.0593351 valid_1's binary_logloss: 0.137471
[10000] training's binary_logloss: 0.0440655 valid_1's binary_logloss: 0.127918
[12000] training's binary_logloss: 0.0340528 valid_1's binary_logloss: 0.124641
[14000] training's binary_logloss: 0.0284739 valid_1's binary_logloss: 0.123799
[16000] training's binary_logloss: 0.0242686 valid_1's binary_logloss: 0.123534
[18000] training's binary_logloss: 0.021996 valid_1's binary_logloss: 0.123676
[20000] training's binary_logloss: 0.0203448 valid_1's binary_logloss: 0.123332
[22000] training's binary_logloss: 0.0190207 valid_1's binary_logloss: 0.117218
[24000] training's binary_logloss: 0.0181458 valid_1's binary_logloss: 0.110904
[26000] training's binary_logloss: 0.0173272 valid_1's binary_logloss: 0.105216
[28000] training's binary_logloss: 0.016914 valid_1's binary_logloss: 0.101382
[30000] training's binary_logloss: 0.0167751 valid_1's binary_logloss: 0.0991921
[32000] training's binary_logloss: 0.0166785 valid_1's binary_logloss: 0.0981959
[34000] training's binary_logloss: 0.0164938 valid_1's binary_logloss: 0.0963396
[36000] training's binary_logloss: 0.0163961 valid_1's binary_logloss: 0.0945048
[38000] training's binary_logloss: 0.0163068 valid_1's binary_logloss: 0.0935821
[40000] training's binary_logloss: 0.0162859 valid_1's binary_logloss: 0.092862
[42000] training's binary_logloss: 0.0162149 valid_1's binary_logloss: 0.0918785
[44000] training's binary_logloss: 0.0160704 valid_1's binary_logloss: 0.0916051
[46000] training's binary_logloss: 0.016084 valid_1's binary_logloss: 0.0909328
[48000] training's binary_logloss: 0.0160163 valid_1's binary_logloss: 0.0905719
[50000] training's binary_logloss: 0.0159573 valid_1's binary_logloss: 0.0902551

lightgbm training fold 3

[2000] training's binary_logloss: 0.245919 valid_1's binary_logloss: 0.359147
[4000] training's binary_logloss: 0.135415 valid_1's binary_logloss: 0.291105
[6000] training's binary_logloss: 0.084923 valid_1's binary_logloss: 0.246192
[8000] training's binary_logloss: 0.0563978 valid_1's binary_logloss: 0.207595
[10000] training's binary_logloss: 0.0425637 valid_1's binary_logloss: 0.186845
[12000] training's binary_logloss: 0.0330554 valid_1's binary_logloss: 0.168603
[14000] training's binary_logloss: 0.0277481 valid_1's binary_logloss: 0.15705
[16000] training's binary_logloss: 0.0237372 valid_1's binary_logloss: 0.149003
[18000] training's binary_logloss: 0.0214382 valid_1's binary_logloss: 0.146831
[20000] training's binary_logloss: 0.0198107 valid_1's binary_logloss: 0.146182
[22000] training's binary_logloss: 0.0185385 valid_1's binary_logloss: 0.146802
[24000] training's binary_logloss: 0.0176257 valid_1's binary_logloss: 0.149218
[26000] training's binary_logloss: 0.016798 valid_1's binary_logloss: 0.150429
[28000] training's binary_logloss: 0.0163229 valid_1's binary_logloss: 0.152989
[30000] training's binary_logloss: 0.0161796 valid_1's binary_logloss: 0.155893
[32000] training's binary_logloss: 0.0161365 valid_1's binary_logloss: 0.159083
[34000] training's binary_logloss: 0.0160855 valid_1's binary_logloss: 0.163047
[36000] training's binary_logloss: 0.0160168 valid_1's binary_logloss: 0.166732
[38000] training's binary_logloss: 0.0159444 valid_1's binary_logloss: 0.169782
[40000] training's binary_logloss: 0.0158308 valid_1's binary_logloss: 0.17239
[42000] training's binary_logloss: 0.0157862 valid_1's binary_logloss: 0.174705
[44000] training's binary_logloss: 0.0157505 valid_1's binary_logloss: 0.177534
[46000] training's binary_logloss: 0.0157089 valid_1's binary_logloss: 0.180396
[48000] training's binary_logloss: 0.0156311 valid_1's binary_logloss: 0.182606
[50000] training's binary_logloss: 0.0156035 valid_1's binary_logloss: 0.18511

lightgbm training fold 4

[2000] training's binary_logloss: 0.241712 valid_1's binary_logloss: 0.354468
[4000] training's binary_logloss: 0.133762 valid_1's binary_logloss: 0.296722
[6000] training's binary_logloss: 0.0822441 valid_1's binary_logloss: 0.285202
[8000] training's binary_logloss: 0.0533863 valid_1's binary_logloss: 0.285733
[10000] training's binary_logloss: 0.0396003 valid_1's binary_logloss: 0.287819
[12000] training's binary_logloss: 0.0306938 valid_1's binary_logloss: 0.289484
[14000] training's binary_logloss: 0.0257471 valid_1's binary_logloss: 0.294007

```
[16000] training's binary_logloss: 0.0220611  valid_1's binary_logloss: 0.297164  
[18000] training's binary_logloss: 0.0200204  valid_1's binary_logloss: 0.3027  
[20000] training's binary_logloss: 0.0186709  valid_1's binary_logloss: 0.310359  
[22000] training's binary_logloss: 0.0175881  valid_1's binary_logloss: 0.313839  
[24000] training's binary_logloss: 0.0168835  valid_1's binary_logloss: 0.318571  
[26000] training's binary_logloss: 0.0161774  valid_1's binary_logloss: 0.324982  
[28000] training's binary_logloss: 0.015865   valid_1's binary_logloss: 0.330759  
[30000] training's binary_logloss: 0.0157936  valid_1's binary_logloss: 0.334225  
[32000] training's binary_logloss: 0.0156188  valid_1's binary_logloss: 0.336284  
[34000] training's binary_logloss: 0.0155631  valid_1's binary_logloss: 0.339592  
[36000] training's binary_logloss: 0.0154838  valid_1's binary_logloss: 0.341127  
[38000] training's binary_logloss: 0.0154582  valid_1's binary_logloss: 0.344447  
[40000] training's binary_logloss: 0.0153956  valid_1's binary_logloss: 0.348196  
[42000] training's binary_logloss: 0.0153609  valid_1's binary_logloss: 0.352446  
[44000] training's binary_logloss: 0.0153071  valid_1's binary_logloss: 0.35542  
[46000] training's binary_logloss: 0.0152281  valid_1's binary_logloss: 0.357728  
[48000] training's binary_logloss: 0.0151642  valid_1's binary_logloss: 0.360626  
[50000] training's binary_logloss: 0.0151072  valid_1's binary_logloss: 0.364738  
-----  
lightgbm training fold 5  
[2000] training's binary_logloss: 0.247366  valid_1's binary_logloss: 0.341665  
[4000] training's binary_logloss: 0.136308  valid_1's binary_logloss: 0.243956  
[6000] training's binary_logloss: 0.083049  valid_1's binary_logloss: 0.215872  
[8000] training's binary_logloss: 0.0541907  valid_1's binary_logloss: 0.206349  
[10000] training's binary_logloss: 0.0405135  valid_1's binary_logloss: 0.200333  
[12000] training's binary_logloss: 0.0314357  valid_1's binary_logloss: 0.196361  
[14000] training's binary_logloss: 0.0266671  valid_1's binary_logloss: 0.19684  
[16000] training's binary_logloss: 0.0229593  valid_1's binary_logloss: 0.196581  
[18000] training's binary_logloss: 0.020887  valid_1's binary_logloss: 0.196581  
[20000] training's binary_logloss: 0.0193373  valid_1's binary_logloss: 0.19775  
[22000] training's binary_logloss: 0.0181064  valid_1's binary_logloss: 0.19707  
[24000] training's binary_logloss: 0.0173531  valid_1's binary_logloss: 0.196517  
[26000] training's binary_logloss: 0.0166049  valid_1's binary_logloss: 0.196913  
[28000] training's binary_logloss: 0.0161905  valid_1's binary_logloss: 0.195339  
[30000] training's binary_logloss: 0.0160866  valid_1's binary_logloss: 0.195122  
[32000] training's binary_logloss: 0.0159147  valid_1's binary_logloss: 0.196548  
[34000] training's binary_logloss: 0.0157572  valid_1's binary_logloss: 0.19736  
[36000] training's binary_logloss: 0.0156425  valid_1's binary_logloss: 0.197567  
[38000] training's binary_logloss: 0.0155601  valid_1's binary_logloss: 0.198249  
[40000] training's binary_logloss: 0.0155023  valid_1's binary_logloss: 0.199145  
[42000] training's binary_logloss: 0.0155052  valid_1's binary_logloss: 0.200517  
[44000] training's binary_logloss: 0.0154495  valid_1's binary_logloss: 0.201886  
[46000] training's binary_logloss: 0.015375  valid_1's binary_logloss: 0.202715  
[48000] training's binary_logloss: 0.0153219  valid_1's binary_logloss: 0.203269  
[50000] training's binary_logloss: 0.0152853  valid_1's binary_logloss: 0.203606  
-----  
lightgbm training fold 6  
[2000] training's binary_logloss: 0.251822  valid_1's binary_logloss: 0.305254  
[4000] training's binary_logloss: 0.137752  valid_1's binary_logloss: 0.2149  
[6000] training's binary_logloss: 0.0841859  valid_1's binary_logloss: 0.178178  
[8000] training's binary_logloss: 0.0550538  valid_1's binary_logloss: 0.164767  
[10000] training's binary_logloss: 0.0411488  valid_1's binary_logloss: 0.157213  
[12000] training's binary_logloss: 0.0319465  valid_1's binary_logloss: 0.151383  
[14000] training's binary_logloss: 0.0267807  valid_1's binary_logloss: 0.152271  
[16000] training's binary_logloss: 0.0229214  valid_1's binary_logloss: 0.151375  
[18000] training's binary_logloss: 0.0208031  valid_1's binary_logloss: 0.15206  
[20000] training's binary_logloss: 0.01932   valid_1's binary_logloss: 0.152681  
[22000] training's binary_logloss: 0.0180624  valid_1's binary_logloss: 0.154805  
[24000] training's binary_logloss: 0.0172254  valid_1's binary_logloss: 0.154982  
[26000] training's binary_logloss: 0.0164263  valid_1's binary_logloss: 0.154331  
[28000] training's binary_logloss: 0.0159945  valid_1's binary_logloss: 0.152144  
[30000] training's binary_logloss: 0.0158889  valid_1's binary_logloss: 0.148446  
[32000] training's binary_logloss: 0.0157587  valid_1's binary_logloss: 0.145524  
[34000] training's binary_logloss: 0.015568   valid_1's binary_logloss: 0.14315  
[36000] training's binary_logloss: 0.0153968  valid_1's binary_logloss: 0.142553  
[38000] training's binary_logloss: 0.0152982  valid_1's binary_logloss: 0.141984  
[40000] training's binary_logloss: 0.0152153  valid_1's binary_logloss: 0.141163  
[42000] training's binary_logloss: 0.0151323  valid_1's binary_logloss: 0.140455  
[44000] training's binary_logloss: 0.0151358  valid_1's binary_logloss: 0.139606  
[46000] training's binary_logloss: 0.0150408  valid_1's binary_logloss: 0.139614  
[48000] training's binary_logloss: 0.0150066  valid_1's binary_logloss: 0.140056  
[50000] training's binary_logloss: 0.0149563  valid_1's binary_logloss: 0.140212  
-----  
lightgbm training fold 7  
[2000] training's binary_logloss: 0.256157  valid_1's binary_logloss: 0.282579  
[4000] training's binary_logloss: 0.14371   valid_1's binary_logloss: 0.179044  
[6000] training's binary_logloss: 0.0892926  valid_1's binary_logloss: 0.135825  
[8000] training's binary_logloss: 0.0587899  valid_1's binary_logloss: 0.11667  
[10000] training's binary_logloss: 0.0439796  valid_1's binary_logloss: 0.106559  
[12000] training's binary_logloss: 0.0339917  valid_1's binary_logloss: 0.0994576  
[14000] training's binary_logloss: 0.028343  valid_1's binary_logloss: 0.0936961  
[16000] training's binary_logloss: 0.0241691  valid_1's binary_logloss: 0.0905641  
[18000] training's binary_logloss: 0.0218713  valid_1's binary_logloss: 0.0906256  
[20000] training's binary_logloss: 0.0201041  valid_1's binary_logloss: 0.0903999  
[22000] training's binary_logloss: 0.018607  valid_1's binary_logloss: 0.0903996  
[24000] training's binary_logloss: 0.0176802  valid_1's binary_logloss: 0.0922518  
[26000] training's binary_logloss: 0.0169163  valid_1's binary_logloss: 0.0939674  
[28000] training's binary_logloss: 0.016554   valid_1's binary_logloss: 0.0944728  
[30000] training's binary_logloss: 0.0163403  valid_1's binary_logloss: 0.0946632  
[32000] training's binary_logloss: 0.0160182  valid_1's binary_logloss: 0.0959939
```

```
[34000] training's binary_logloss: 0.015841          valid_1's binary_logloss: 0.0979622
[36000] training's binary_logloss: 0.0158555         valid_1's binary_logloss: 0.0995334
[38000] training's binary_logloss: 0.0158898         valid_1's binary_logloss: 0.101227
[40000] training's binary_logloss: 0.0158912         valid_1's binary_logloss: 0.103071
[42000] training's binary_logloss: 0.0159894         valid_1's binary_logloss: 0.104938
[44000] training's binary_logloss: 0.0159996         valid_1's binary_logloss: 0.106075
[46000] training's binary_logloss: 0.0159974         valid_1's binary_logloss: 0.107834
[48000] training's binary_logloss: 0.0160343         valid_1's binary_logloss: 0.109342
[50000] training's binary_logloss: 0.0160203         valid_1's binary_logloss: 0.11061
-----
lightgbm training fold 8
[2000] training's binary_logloss: 0.249766          valid_1's binary_logloss: 0.317574
[4000] training's binary_logloss: 0.139762          valid_1's binary_logloss: 0.231413
[6000] training's binary_logloss: 0.087177          valid_1's binary_logloss: 0.193275
[8000] training's binary_logloss: 0.0570707         valid_1's binary_logloss: 0.169029
[10000] training's binary_logloss: 0.0423716         valid_1's binary_logloss: 0.155111
[12000] training's binary_logloss: 0.0325945         valid_1's binary_logloss: 0.146866
[14000] training's binary_logloss: 0.0273896         valid_1's binary_logloss: 0.141895
[16000] training's binary_logloss: 0.0234877         valid_1's binary_logloss: 0.139626
[18000] training's binary_logloss: 0.0213089         valid_1's binary_logloss: 0.138133
[20000] training's binary_logloss: 0.0197246         valid_1's binary_logloss: 0.136997
[22000] training's binary_logloss: 0.0183928         valid_1's binary_logloss: 0.135534
[24000] training's binary_logloss: 0.0175528         valid_1's binary_logloss: 0.133181
[26000] training's binary_logloss: 0.0168514         valid_1's binary_logloss: 0.132824
[28000] training's binary_logloss: 0.0164903         valid_1's binary_logloss: 0.132744
[30000] training's binary_logloss: 0.0163053         valid_1's binary_logloss: 0.132812
[32000] training's binary_logloss: 0.016056          valid_1's binary_logloss: 0.132853
[34000] training's binary_logloss: 0.0159538         valid_1's binary_logloss: 0.132945
[36000] training's binary_logloss: 0.0158875         valid_1's binary_logloss: 0.133574
[38000] training's binary_logloss: 0.0158548         valid_1's binary_logloss: 0.134527
[40000] training's binary_logloss: 0.0157682         valid_1's binary_logloss: 0.13497
[42000] training's binary_logloss: 0.0157762         valid_1's binary_logloss: 0.135267
[44000] training's binary_logloss: 0.0157078         valid_1's binary_logloss: 0.135402
[46000] training's binary_logloss: 0.0156055         valid_1's binary_logloss: 0.134848
[48000] training's binary_logloss: 0.0155719         valid_1's binary_logloss: 0.134877
[50000] training's binary_logloss: 0.0154665         valid_1's binary_logloss: 0.134386
-----
lightgbm training fold 9
[2000] training's binary_logloss: 0.249132          valid_1's binary_logloss: 0.330581
[4000] training's binary_logloss: 0.13719          valid_1's binary_logloss: 0.259868
[6000] training's binary_logloss: 0.0856526          valid_1's binary_logloss: 0.247146
[8000] training's binary_logloss: 0.0564939          valid_1's binary_logloss: 0.242483
[10000] training's binary_logloss: 0.0421256          valid_1's binary_logloss: 0.243425
[12000] training's binary_logloss: 0.032474          valid_1's binary_logloss: 0.247867
[14000] training's binary_logloss: 0.0270399          valid_1's binary_logloss: 0.254305
[16000] training's binary_logloss: 0.0230066          valid_1's binary_logloss: 0.258114
[18000] training's binary_logloss: 0.020673          valid_1's binary_logloss: 0.260256
[20000] training's binary_logloss: 0.0190398          valid_1's binary_logloss: 0.255699
[22000] training's binary_logloss: 0.017727          valid_1's binary_logloss: 0.257821
[24000] training's binary_logloss: 0.0168632          valid_1's binary_logloss: 0.26109
[26000] training's binary_logloss: 0.0160323          valid_1's binary_logloss: 0.264075
[28000] training's binary_logloss: 0.0155176          valid_1's binary_logloss: 0.266002
[30000] training's binary_logloss: 0.0152811          valid_1's binary_logloss: 0.267597
[32000] training's binary_logloss: 0.0150165          valid_1's binary_logloss: 0.270536
[34000] training's binary_logloss: 0.014719          valid_1's binary_logloss: 0.274964
[36000] training's binary_logloss: 0.0144499          valid_1's binary_logloss: 0.280472
[38000] training's binary_logloss: 0.0143374          valid_1's binary_logloss: 0.281893
[40000] training's binary_logloss: 0.0142739          valid_1's binary_logloss: 0.283633
[42000] training's binary_logloss: 0.0142475          valid_1's binary_logloss: 0.285381
[44000] training's binary_logloss: 0.0142173          valid_1's binary_logloss: 0.286783
[46000] training's binary_logloss: 0.0141486          valid_1's binary_logloss: 0.288023
[48000] training's binary_logloss: 0.0141513          valid_1's binary_logloss: 0.290197
[50000] training's binary_logloss: 0.0140934          valid_1's binary_logloss: 0.292563
-----
lightgbm training fold 10
[2000] training's binary_logloss: 0.257802          valid_1's binary_logloss: 0.304896
[4000] training's binary_logloss: 0.143402          valid_1's binary_logloss: 0.217758
[6000] training's binary_logloss: 0.0877063          valid_1's binary_logloss: 0.179941
[8000] training's binary_logloss: 0.0564943          valid_1's binary_logloss: 0.15912
[10000] training's binary_logloss: 0.0415212          valid_1's binary_logloss: 0.148131
[12000] training's binary_logloss: 0.0317468          valid_1's binary_logloss: 0.1419
[14000] training's binary_logloss: 0.0265053          valid_1's binary_logloss: 0.137396
[16000] training's binary_logloss: 0.0226465          valid_1's binary_logloss: 0.134378
[18000] training's binary_logloss: 0.0204782          valid_1's binary_logloss: 0.132486
[20000] training's binary_logloss: 0.0189869          valid_1's binary_logloss: 0.1304
[22000] training's binary_logloss: 0.0178852          valid_1's binary_logloss: 0.129018
[24000] training's binary_logloss: 0.0171111          valid_1's binary_logloss: 0.128068
[26000] training's binary_logloss: 0.0163095          valid_1's binary_logloss: 0.128016
[28000] training's binary_logloss: 0.0158448          valid_1's binary_logloss: 0.127429
[30000] training's binary_logloss: 0.0157027          valid_1's binary_logloss: 0.126766
[32000] training's binary_logloss: 0.0154712          valid_1's binary_logloss: 0.126664
[34000] training's binary_logloss: 0.015415          valid_1's binary_logloss: 0.126957
[36000] training's binary_logloss: 0.0152814          valid_1's binary_logloss: 0.127199
[38000] training's binary_logloss: 0.0152584          valid_1's binary_logloss: 0.127962
[40000] training's binary_logloss: 0.0152167          valid_1's binary_logloss: 0.128843
[42000] training's binary_logloss: 0.015232          valid_1's binary_logloss: 0.129459
[44000] training's binary_logloss: 0.0152039          valid_1's binary_logloss: 0.129409
[46000] training's binary_logloss: 0.0151354          valid_1's binary_logloss: 0.12968
[48000] training's binary_logloss: 0.0150907          valid_1's binary_logloss: 0.130332
[50000] training's binary_logloss: 0.0150543          valid_1's binary_logloss: 0.130257
```

```
lightgbm our out of folds CV score is 0.2433478814875392
-----
xgboost training fold 1
[0]    train-logloss:0.68955  eval-logloss:0.69139
[1049]  train-logloss:0.04785  eval-logloss:0.42206
-----
xgboost training fold 2
[0]    train-logloss:0.68980  eval-logloss:0.69074
[2000]  train-logloss:0.02634  eval-logloss:0.16347
[4000]  train-logloss:0.02197  eval-logloss:0.15946
[4021]  train-logloss:0.02195  eval-logloss:0.15945
-----
xgboost training fold 3
[0]    train-logloss:0.68966  eval-logloss:0.69015
[2000]  train-logloss:0.02572  eval-logloss:0.28721
[2641]  train-logloss:0.02333  eval-logloss:0.28584
-----
xgboost training fold 4
[0]    train-logloss:0.68948  eval-logloss:0.68999
[1056]  train-logloss:0.04640  eval-logloss:0.42808
-----
xgboost training fold 5
[0]    train-logloss:0.68972  eval-logloss:0.69013
[2000]  train-logloss:0.02536  eval-logloss:0.25311
[3652]  train-logloss:0.02184  eval-logloss:0.24964
-----
xgboost training fold 6
[0]    train-logloss:0.68961  eval-logloss:0.69091
[2000]  train-logloss:0.02610  eval-logloss:0.31762
[2141]  train-logloss:0.02534  eval-logloss:0.31758
-----
xgboost training fold 7
[0]    train-logloss:0.68969  eval-logloss:0.68963
[2000]  train-logloss:0.02639  eval-logloss:0.18459
[4000]  train-logloss:0.02192  eval-logloss:0.17804
[4525]  train-logloss:0.02150  eval-logloss:0.17832
-----
xgboost training fold 8
[0]    train-logloss:0.68965  eval-logloss:0.69105
[2000]  train-logloss:0.02540  eval-logloss:0.16245
[4000]  train-logloss:0.02140  eval-logloss:0.15699
[4242]  train-logloss:0.02125  eval-logloss:0.15714
-----
xgboost training fold 9
[0]    train-logloss:0.68962  eval-logloss:0.69116
[1516]  train-logloss:0.03121  eval-logloss:0.31140
-----
xgboost training fold 10
[0]   train-logloss:0.68953  eval-logloss:0.68989
[2000]  train-logloss:0.02554  eval-logloss:0.13727
[4000]  train-logloss:0.02116  eval-logloss:0.13203
[6000]  train-logloss:0.02009  eval-logloss:0.12967
[8000]  train-logloss:0.01945  eval-logloss:0.12838
[10000] train-logloss:0.01907  eval-logloss:0.12795
[10012] train-logloss:0.01905  eval-logloss:0.12795
xgboost our out of folds CV score is 0.36168974268250453
-----
catboost training fold 1
0:    learn: 0.6899332      test: 0.6907982 best: 0.6907982 (0)      total: 58.1ms  remaining: 48m 54s
2000:  learn: 0.0414667     test: 0.3914124 best: 0.3802861 (1519)  total: 3.41s  remaining: 1m 22s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.3802861273
bestIteration = 1519

Shrink model to first 1520 iterations.
-----
catboost training fold 2
0:    learn: 0.6902880      test: 0.6897226 best: 0.6897226 (0)      total: 1.88ms  remaining: 1m 35s
2000:  learn: 0.0455178     test: 0.1893871 best: 0.1890963 (1983)  total: 3.3s  remaining: 1m 20s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.174910968
bestIteration = 2391

Shrink model to first 2392 iterations.
-----
catboost training fold 3
0:    learn: 0.6899973      test: 0.6903132 best: 0.6903132 (0)      total: 1.7ms  remaining: 1m 26s
2000:  learn: 0.0454009     test: 0.2314102 best: 0.2310818 (1996)  total: 3.31s  remaining: 1m 20s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.2227033843
bestIteration = 2594

Shrink model to first 2595 iterations.
-----
catboost training fold 4
0:    learn: 0.6902377      test: 0.6903190 best: 0.6903190 (0)      total: 1.8ms  remaining: 1m 30s
Stopped by overfitting detector (500 iterations wait)
```

```

bestTest = 0.3824828654
bestIteration = 1023

Shrink model to first 1024 iterations.
-----
catboost training fold 5
0:    learn: 0.6902422      test: 0.6899385 best: 0.6899385 (0)      total: 1.97ms  remaining: 1m 39s
2000:   learn: 0.0419670      test: 0.2187739 best: 0.2187590 (1995)  total: 3.28s  remaining: 1m 19s
4000:   learn: 0.0081315      test: 0.1913288 best: 0.1913245 (3989)  total: 6.66s  remaining: 1m 17s
6000:   learn: 0.0037223      test: 0.1836002 best: 0.1834333 (5990)  total: 9.91s  remaining: 1m 13s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.1828987159
bestIteration = 6258

Shrink model to first 6259 iterations.
-----
catboost training fold 6
0:    learn: 0.6902818      test: 0.6906277 best: 0.6906277 (0)      total: 1.71ms  remaining: 1m 26s
2000:   learn: 0.0465258      test: 0.2728627 best: 0.2702758 (1706)  total: 3.38s  remaining: 1m 21s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.2702757782
bestIteration = 1706

Shrink model to first 1707 iterations.
-----
catboost training fold 7
0:    learn: 0.6901248      test: 0.6907622 best: 0.6907622 (0)      total: 1.95ms  remaining: 1m 38s
2000:   learn: 0.0448237      test: 0.1739162 best: 0.1738803 (1995)  total: 3.33s  remaining: 1m 20s
4000:   learn: 0.0091614      test: 0.1518411 best: 0.1518221 (3999)  total: 6.63s  remaining: 1m 17s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.1474822276
bestIteration = 4828

Shrink model to first 4829 iterations.
-----
catboost training fold 8
0:    learn: 0.6906348      test: 0.6923224 best: 0.6923224 (0)      total: 2.02ms  remaining: 1m 41s
2000:   learn: 0.0459817      test: 0.1827642 best: 0.1827270 (1999)  total: 3.59s  remaining: 1m 27s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.1636855918
bestIteration = 2864

Shrink model to first 2865 iterations.
-----
catboost training fold 9
0:    learn: 0.6892691      test: 0.6898655 best: 0.6898655 (0)      total: 1.79ms  remaining: 1m 30s
2000:   learn: 0.0461002      test: 0.3204117 best: 0.3190725 (1908)  total: 3.31s  remaining: 1m 20s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.3190725037
bestIteration = 1908

Shrink model to first 1909 iterations.
-----
catboost training fold 10
0:    learn: 0.6908658      test: 0.6907668 best: 0.6907668 (0)      total: 1.91ms  remaining: 1m 36s
2000:   learn: 0.0460972      test: 0.1584097 best: 0.1583195 (1999)  total: 3.29s  remaining: 1m 19s
4000:   learn: 0.0091052      test: 0.1137035 best: 0.1136661 (3957)  total: 6.56s  remaining: 1m 16s
Stopped by overfitting detector (500 iterations wait)

bestTest = 0.1096450309
bestIteration = 5000

Shrink model to first 5001 iterations.
catboost our out of folds CV score is 0.3166778761418612

```

Inference

```

In [21]: def lightgbm_inference(x_test: pd.DataFrame):
    test_pred = np.zeros(len(x_test))
    for fold in range(CFG.n_folds):
        model = pickle.load(open(CFG.MODEL_DATA_PATH / f'lightgbm_fold{fold + 1}_seed{CFG.seed}_ver{CFG.VER}.pkl', 'rb'))
        # Predict
        test_pred += model.predict(x_test)
    return test_pred / CFG.n_folds

def xgboost_inference(x_test: pd.DataFrame):
    test_pred = np.zeros(len(x_test))
    for fold in range(CFG.n_folds):
        model = pickle.load(open(CFG.MODEL_DATA_PATH / f'xgboost_fold{fold + 1}_seed{CFG.seed}_ver{CFG.VER}.pkl', 'rb'))
        # Predict
        test_pred += model.predict(xgb.DMatrix(x_test), iteration_range=(0, model.best_ntree_limit))
    return test_pred / CFG.n_folds

def catboost_inference(x_test: pd.DataFrame):
    test_pred = np.zeros(len(x_test))

```

```

for fold in range(CFG.n_folds):
    model = pickle.load(open(CFG.MODEL_DATA_PATH / f'catboost_fold{fold + 1}_seed{CFG.seed}_ver{CFG.VER}.pkl', 'rb'))
    # Predict
    test_pred += model.predict_proba(x_test)[:, 1]
return test_pred / CFG.n_folds

def gradient_boosting_model_inference(method: str, test_df: pd.DataFrame, features: list, categorical_features: list):
    x_test = test_df[features]
    if method == 'lightgbm':
        test_pred = lightgbm_inference(x_test)
    if method == 'xgboost':
        test_pred = xgboost_inference(x_test)
    if method == 'catboost':
        test_pred = catboost_inference(x_test)
    return test_pred

```

In [22]: `for method in CFG.METHOD_LIST:
 test_df[f'{method}_pred_prob'] = gradient_boosting_model_inference(method, test_df, features, categorical_features)`

The result

In [23]: `test_df['class_1'] = 0.9 * test_df['lightgbm_pred_prob'] + 0.05 * test_df['xgboost_pred_prob'] + 0.05 * test_df['catboost_pred_prob']
test_df['class_0'] = 1 - test_df['class_1']
test_df[list(result_df)].to_csv('result.csv', index=False)`

In [24]: `test_df[list(result_df)]`

Out[24]:

	Id	class_0	class_1
0	00eed32682bb	0.641967	0.358033
1	010ebe33f668	0.641967	0.358033
2	02fa521e1838	0.641967	0.358033
3	040e15f562a2	0.641967	0.358033
4	046e85c7cc7f	0.641967	0.358033